

# EDA\_Springleaf\_screencast

December 23, 2019

This is a notebook, used in the screencast video. Note, that the data files are not present here in Jupyter hub and you will not be able to run it. But you can always download the notebook to your local machine as well as the competition data and make it interactive.

```
In [1]: import os
import numpy as np
import pandas as pd
from tqdm import tqdm_notebook
import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')

import seaborn

In [2]: def autolabel(arrayA):
    ''' label each colored square with the corresponding data value.
    If value > 20, the text is in black, else in white.
    '''
    arrayA = np.array(arrayA)
    for i in range(arrayA.shape[0]):
        for j in range(arrayA.shape[1]):
            plt.text(j,i, "%.2f"%arrayA[i,j], ha='center', va='bottom',color='w')

def hist_it(feats):
    plt.figure(figsize=(16,4))
    feat[Y==0].hist(bins=range(int(feats.min()),int(feats.max()+2)),normed=True,alpha=0.8)
    feat[Y==1].hist(bins=range(int(feats.min()),int(feats.max()+2)),normed=True,alpha=0.5)
    plt.ylim((0,1))

def gt_matrix(feats,sz=16):
    a = []
    for i,c1 in enumerate(feats):
        b = []
        for j,c2 in enumerate(feats):
            mask = (~train[c1].isnull()) & (~train[c2].isnull())
```

```

        if i>=j:
            b.append((train.loc[mask,c1].values>=train.loc[mask,c2].values).mean())
        else:
            b.append((train.loc[mask,c1].values>train.loc[mask,c2].values).mean())

    a.append(b)

plt.figure(figsize = (sz,sz))
plt.imshow(a, interpolation = 'None')
_ = plt.xticks(range(len(feats)),feats,rotation = 90)
_ = plt.yticks(range(len(feats)),feats,rotation = 0)
autolabel(a)

In [3]: def hist_it1(feat):
    plt.figure(figsize=(16,4))
    feat[Y==0].hist(bins=100,range=(feat.min(),feat.max()),normed=True,alpha=0.5)
    feat[Y==1].hist(bins=100,range=(feat.min(),feat.max()),normed=True,alpha=0.5)
    plt.ylim((0,1))

```

## 1 Read the data

```

In [4]: train = pd.read_csv('train.csv.zip')
        Y = train.target

In [5]: test = pd.read_csv('test.csv.zip')
        test_ID = test.ID

```

## 2 Data overview

Probably the first thing you check is the shapes of the train and test matrices and look inside them.

```

In [6]: print 'Train shape', train.shape
        print 'Test shape', test.shape

```

```

Train shape (145231, 1934)
Test shape (145232, 1933)

```

```

In [10]: train.head()

```

```

Out[10]:
   ID  VAR_0001  VAR_0002  VAR_0003  VAR_0004  VAR_0005  VAR_0006  VAR_0007  \
0    2         H        224          0       4300         C         0.0         0.0
1    4         H         7         53       4448         B         1.0         0.0
2    5         H        116          3       3464         C         0.0         0.0
3    7         H        240        300       3200         C         0.0         0.0
4    8         R         72        261       2000         N         0.0         0.0

   VAR_0008  VAR_0009  ...  VAR_1926  VAR_1927  VAR_1928  VAR_1929  VAR_1930  \

```

0	False	False	...	98	98	998	999999998	998
1	False	False	...	98	98	998	999999998	998
2	False	False	...	98	98	998	999999998	998
3	False	False	...	98	98	998	999999998	998
4	False	False	...	98	98	998	999999998	998

	VAR_1931	VAR_1932	VAR_1933	VAR_1934	target
0	998	9998	9998	IAPS	0
1	998	9998	9998	IAPS	0
2	998	9998	9998	IAPS	0
3	998	9998	9998	RCC	0
4	998	9998	9998	BRANCH	1

[5 rows x 1934 columns]

In [11]: test.head()

```
Out[11]:
```

	ID	VAR_0001	VAR_0002	VAR_0003	VAR_0004	VAR_0005	VAR_0006	VAR_0007	\
0	1	R	360	25	2251	B	2.0	2.0	
1	3	R	74	192	3274	C	2.0	3.0	
2	6	R	21	36	3500	C	1.0	1.0	
3	9	R	8	2	1500	B	0.0	0.0	
4	10	H	91	39	84500	C	8.0	3.0	

	VAR_0008	VAR_0009	...	VAR_1925	VAR_1926	VAR_1927	VAR_1928	VAR_1929	\
0	False	False	...	0	98	98	998	999999998	
1	False	False	...	0	98	98	998	999999998	
2	False	False	...	0	98	98	998	999999998	
3	False	False	...	0	98	98	998	999999998	
4	False	False	...	0	98	98	998	999999998	

	VAR_1930	VAR_1931	VAR_1932	VAR_1933	VAR_1934
0	998	998	9998	9998	IAPS
1	998	998	9998	9998	IAPS
2	998	998	9998	9998	IAPS
3	998	998	9998	9998	IAPS
4	998	998	9998	9998	IAPS

[5 rows x 1933 columns]

There are almost 2000 anonymized variables! It's clear, some of them are categorical, some look like numeric. Some numeric features are integer typed, so probably they are event counters or dates. And others are of float type, but from the first few rows they look like integer-typed too, since fractional part is zero, but pandas treats them as float since there are NaN values in that features.

From the first glance we see train has one more column target which we should not forget to drop before fitting a classifier. We also see ID column is shared between train and test, which sometimes can be successfully used to improve the score.

It is also useful to know if there are any NaNs in the data. You should pay attention to columns with NaNs and the number of NaNs for each row can serve as a nice feature later.

```
In [13]: # Number of NaNs for each object
         train.isnull().sum(axis=1).head(15)
```

```
Out[13]: 0      25
         1      19
         2      24
         3      24
         4      24
         5      24
         6      24
         7      24
         8      16
         9      24
        10      22
        11      24
        12      17
        13      24
        14      24
         dtype: int64
```

```
In [14]: # Number of NaNs for each column
         train.isnull().sum(axis=0).head(15)
```

```
Out[14]: ID      0
         VAR_0001  0
         VAR_0002  0
         VAR_0003  0
         VAR_0004  0
         VAR_0005  0
         VAR_0006  56
         VAR_0007  56
         VAR_0008  56
         VAR_0009  56
         VAR_0010  56
         VAR_0011  56
         VAR_0012  56
         VAR_0013  56
         VAR_0014  56
         dtype: int64
```

Just by reviewing the head of the lists we immediately see the patterns, exactly 56 NaNs for a set of variables, and 24 NaNs for objects.

## 3 Dataset cleaning

### 3.0.1 Remove constant features

All 1932 columns are anonymized which makes us to deduce the meaning of the features ourselves. We will now try to clean the dataset.

It is usually convenient to concatenate train and test into one dataframe and do all feature engineering using it.

```
In [15]: traintest = pd.concat([train, test], axis = 0)
```

First we should look for a constant features, such features do not provide any information and only make our dataset larger.

```
In [16]: # `dropna = False` makes unique treat NaNs as a distinct value
        feats_counts = train.nunique(dropna = False)
```

```
In [17]: feats_counts.sort_values()[:10]
```

```
Out[17]: VAR_0213      1
         VAR_0207      1
         VAR_0840      1
         VAR_0847      1
         VAR_1428      1
         VAR_1165      2
         VAR_0438      2
         VAR_1164      2
         VAR_1163      2
         VAR_1162      2
         dtype: int64
```

We found 5 constant features. Let's remove them.

```
In [18]: constant_features = feats_counts.loc[feats_counts==1].index.tolist()
        print (constant_features)
```

```
traintest.drop(constant_features,axis = 1,inplace=True)
```

```
['VAR_0207', 'VAR_0213', 'VAR_0840', 'VAR_0847', 'VAR_1428']
```

### 3.0.2 Remove duplicated features

Fill NaNs with something we can find later if needed.

```
In [19]: traintest.fillna('NaN', inplace=True)
```

Now let's encode each feature, as we discussed.

```
In [32]: train_enc = pd.DataFrame(index = train.index)
```

```
    for col in tqdm_notebook(traintest.columns):  
        train_enc[col] = train[col].factorize()[0]
```

We could also do something like this:

```
In [33]: # train_enc[col] = train[col].map(train[col].value_counts())
```

The resulting data frame is very very large, so we cannot just transpose it and use `.duplicated`. That is why we will use a simple loop.

```
In [34]: dup_cols = {}
```

```
    for i, c1 in enumerate(tqdm_notebook(train_enc.columns)):  
        for c2 in train_enc.columns[i + 1:]:  
            if c2 not in dup_cols and np.all(train_enc[c1] == train_enc[c2]):  
                dup_cols[c2] = c1
```

```
In [36]: dup_cols
```

```
Out[36]: {'VAR_0009': 'VAR_0008',  
          'VAR_0010': 'VAR_0008',  
          'VAR_0011': 'VAR_0008',  
          'VAR_0012': 'VAR_0008',  
          'VAR_0013': 'VAR_0006',  
          'VAR_0018': 'VAR_0008',  
          'VAR_0019': 'VAR_0008',  
          'VAR_0020': 'VAR_0008',  
          'VAR_0021': 'VAR_0008',  
          'VAR_0022': 'VAR_0008',  
          'VAR_0023': 'VAR_0008',  
          'VAR_0024': 'VAR_0008',  
          'VAR_0025': 'VAR_0008',  
          'VAR_0026': 'VAR_0008',  
          'VAR_0027': 'VAR_0008',  
          'VAR_0028': 'VAR_0008',  
          'VAR_0029': 'VAR_0008',  
          'VAR_0030': 'VAR_0008',  
          'VAR_0031': 'VAR_0008',  
          'VAR_0032': 'VAR_0008',  
          'VAR_0038': 'VAR_0008',  
          'VAR_0039': 'VAR_0008',
```

```

'VAR_0040': 'VAR_0008',
'VAR_0041': 'VAR_0008',
'VAR_0042': 'VAR_0008',
'VAR_0043': 'VAR_0008',
'VAR_0044': 'VAR_0008',
'VAR_0181': 'VAR_0180',
'VAR_0182': 'VAR_0180',
'VAR_0189': 'VAR_0188',
'VAR_0190': 'VAR_0188',
'VAR_0196': 'VAR_0008',
'VAR_0197': 'VAR_0008',
'VAR_0199': 'VAR_0008',
'VAR_0201': 'VAR_0051',
'VAR_0202': 'VAR_0008',
'VAR_0203': 'VAR_0008',
'VAR_0210': 'VAR_0208',
'VAR_0211': 'VAR_0208',
'VAR_0215': 'VAR_0008',
'VAR_0216': 'VAR_0008',
'VAR_0221': 'VAR_0008',
'VAR_0222': 'VAR_0008',
'VAR_0223': 'VAR_0008',
'VAR_0228': 'VAR_0227',
'VAR_0229': 'VAR_0008',
'VAR_0238': 'VAR_0089',
'VAR_0239': 'VAR_0008',
'VAR_0357': 'VAR_0260',
'VAR_0394': 'VAR_0246',
'VAR_0438': 'VAR_0246',
'VAR_0446': 'VAR_0246',
'VAR_0512': 'VAR_0506',
'VAR_0527': 'VAR_0246',
'VAR_0528': 'VAR_0246',
'VAR_0529': 'VAR_0526',
'VAR_0530': 'VAR_0246',
'VAR_0672': 'VAR_0670',
'VAR_1036': 'VAR_0916'}

```

Don't forget to save them, as it takes long time to find these.

```

In [37]: import cPickle as pickle
         pickle.dump(dup_cols, open('dup_cols.p', 'w'), protocol=pickle.HIGHEST_PROTOCOL)

```

Drop from traintest.

```

In [38]: traintest.drop(dup_cols.keys(), axis = 1,inplace=True)

```

## 4 Determine types

Let's examine the number of unique values.

```
In [50]: nunique = train.nunique(dropna=False)
         nunique
```


```
Out[50]: ID          145231
         VAR_0001         3
         VAR_0002        820
         VAR_0003        588
         VAR_0004       7935
         VAR_0005         4
         VAR_0006        38
         VAR_0007        36
         VAR_0008         2
         VAR_0009         2
         VAR_0010         2
         VAR_0011         2
         VAR_0012         2
         VAR_0013        38
         VAR_0014        38
         VAR_0015        27
         VAR_0016        30
         VAR_0017        26
         VAR_0018         2
         VAR_0019         2
         VAR_0020         2
         VAR_0021         2
         VAR_0022         2
         VAR_0023         2
         VAR_0024         2
         VAR_0025         2
         VAR_0026         2
         VAR_0027         2
         VAR_0028         2
         VAR_0029         2
         ...
         VAR_1907         41
         VAR_1908         37
         VAR_1909         41
         VAR_1910         37
         VAR_1911        107
         VAR_1912       16370
         VAR_1913       25426
         VAR_1914       14226
         VAR_1915        1148
         VAR_1916         8
```



VAR_1917	10
VAR_1918	86
VAR_1919	383
VAR_1920	22
VAR_1921	18
VAR_1922	6798
VAR_1923	2445
VAR_1924	573
VAR_1925	11
VAR_1926	6
VAR_1927	10
VAR_1928	30
VAR_1929	591
VAR_1930	8
VAR_1931	10
VAR_1932	74
VAR_1933	363
VAR_1934	5
target	2
VAR_0004_mod50	50
Length: 1935, dtype: int64	

and build a histogram of those values

```
In [44]: plt.figure(figsize=(14,6))
         _ = plt.hist(nunique.astype(float)/train.shape[0], bins=100)
```



Let's take a look at the features with a huge number of unique values:

```
In [61]: mask = (nunique.astype(float)/train.shape[0] > 0.8)
         train.loc[:, mask]
```

```

Out[61]:
      ID      VAR_0212      VAR_0227
0         2          NaN        311951
1         4  9.20713e+10  2.76949e+06
2         5  2.65477e+10        654127
3         7  7.75753e+10  3.01509e+06
4         8  6.04238e+10        118678
5        14  7.73796e+10  1.76557e+06
6        16  9.70303e+10        80151
7        20  3.10981e+10        853641
8        21  7.82124e+10  1.40254e+06
9        22  1.94014e+10  2.2187e+06
10       23  3.71295e+10  2.77679e+06
11       24  3.01203e+10        434300
12       25  1.80185e+10  1.48914e+06
13       26  9.83358e+10        686666
14       28  9.33087e+10  1.4847e+06
15       30  2.01715e+10        883714
16       31  4.15638e+10  2.6707e+06
17       32  9.17617e+10  2.65485e+06
18       35  3.81344e+10        487721
19       36          NaN  2.54705e+06
20       37  3.27144e+10  1.74684e+06
21       38  1.82142e+10  2.5813e+06
22       40  7.70153e+10  2.59396e+06
23       42  4.69701e+10  1.02977e+06
24       43  9.84442e+10  1.45101e+06
25       46          NaN  2.37136e+06
26       50  9.25094e+10        665930
27       51  3.09094e+10        497686
28       52  6.06105e+10  1.95816e+06
29       54  3.78768e+10  1.62591e+06
...      ...      ...      ...
145201  290409  8.80126e+10  1.83053e+06
145202  290412  4.6152e+10  1.02024e+06
145203  290414  9.33055e+10  1.88151e+06
145204  290415  4.63509e+10        669351
145205  290417  2.36028e+10        655797
145206  290424  3.73293e+10  1.45626e+06
145207  290426  2.38892e+10  1.9503e+06
145208  290427  6.38632e+10        596365
145209  290429  3.00602e+10        572119
145210  290431  4.33429e+10        16120
145211  290432  3.86543e+10  2.08375e+06
145212  290434  9.21391e+10  1.89779e+06
145213  290436  3.07472e+10  2.94532e+06
145214  290439  7.83326e+10  2.54726e+06
145215  290440          NaN        600318
145216  290441  2.78561e+10        602505

```

145217	290443	1.90952e+10	2.44184e+06
145218	290445	4.62035e+10	2.87349e+06
145219	290447	NaN	1.53493e+06
145220	290448	7.54282e+10	1.60102e+06
145221	290449	4.30768e+10	2.08415e+06
145222	290450	7.81325e+10	2.85367e+06
145223	290452	4.51061e+10	1.56506e+06
145224	290453	4.62223e+10	1.46815e+06
145225	290454	7.74507e+10	2.92811e+06
145226	290457	7.05088e+10	2.03657e+06
145227	290458	9.02492e+10	1.68013e+06
145228	290459	9.17224e+10	2.41922e+06
145229	290461	4.51033e+10	1.53960e+06
145230	290463	9.14114e+10	2.6609e+06

[145231 rows x 3 columns]

The values are not float, they are integer, so these features are likely to be even counts. Let's look at another pack of features.

```
In [64]: mask = (nunique.astype(float)/train.shape[0] < 0.8) & (nunique.astype(float)/train.shape[0] < 0.8)
train.loc[:, mask]
```

```
Out[64]:
```

	VAR_0541	VAR_0543	VAR_0899	VAR_1081	VAR_1082	VAR_1087	\
0	49463	116783	112871	76857	76857	116783	
1	303472	346196	346375	341365	341365	346196	
2	94990	122601	121501	107267	107267	121501	
3	20593	59490	61890	45794	47568	59490	
4	10071	35708	34787	20475	23647	34708	
5	18877	28055	28455	21139	21139	28055	
6	321783	333565	886886	327744	327744	333565	
7	2961	5181	11084	4326	4326	5181	
8	20359	30114	33434	24969	27128	30114	
9	815	1300	7677	1197	1197	1300	
10	6088	15233	15483	7077	7077	15233	
11	432	1457	2000	621	621	757	
12	383	539	860	752	1158	539	
13	14359	47562	47562	17706	17706	47562	
14	145391	218067	214836	176627	176627	216307	
15	10040	12119	17263	10399	10399	12119	
16	4880	9607	9607	9165	9165	9607	
17	12900	35590	35781	26096	26096	35590	
18	104442	139605	150505	136419	142218	139605	
19	13898	25566	26685	20122	20122	25566	
20	3524	10033	10133	5838	5838	10033	
21	129873	204072	206946	183049	183049	204072	
22	3591	11400	17680	5565	5565	11400	
23	999999999	999999999	-99999	999999999	999999999	999999999	

24	1270	4955	12201	2490	2490	4955
25	2015	2458	2458	2015	2015	2458
	VAR_1179	VAR_1180	VAR_1181			
0	76857	76857	76857			
1	341365	341365	176604			
2	107267	107267	58714			
3	45794	47568	47568			
4	20475	23647	23647			
5	21139	21139	20627			
6	327744	327744	163944			
7	4326	4326	4326			
8	24969	27128	27128			
9	1197	1197	1197			
10	7077	7077	4033			
11	621	621	621			
12	752	1158	1158			
13	17706	17706	17706			
14	175273	175273	91019			
15	10399	10399	5379			
16	9165	9165	9165			
17	26096	26096	19646			
18	136419	142218	142218			
19	20122	20122	20122			
20	5838	5838	5838			
21	183049	183049	96736			
22	5565	5565	5565			
23	999999999	999999999	999999999			
24	2490	2490	2490			
25	2015	2015	1008			

These look like counts too. First thing to notice is the 23th line: 99999., -99999 values look like NaNs so we should probably built a related feature. Second: the columns are sometimes placed next to each other, so the columns are probably grouped together and we can disentangle that.

Our conclusion: there are no floating point variables, there are some counts variables, which we will treat as numeric.

And finally, let's pick one variable (in this case 'VAR\_0015') from the third group of features.

```
In [82]: train['VAR_0015'].value_counts()
```

```
Out[82]: 0.0    102382
         1.0    28045
         2.0     8981
         3.0    3199
         4.0    1274
         5.0     588
         6.0     275
         7.0     166
```

8.0	97
-999.0	56
9.0	51
10.0	39
11.0	18
12.0	16
13.0	9
14.0	8
15.0	8
16.0	6
22.0	3
21.0	3
19.0	1
35.0	1
17.0	1
29.0	1
18.0	1
32.0	1
23.0	1

Name: VAR\_0015, dtype: int64

```
In [5]: cat_cols = list(train.select_dtypes(include=['object']).columns)
        num_cols = list(train.select_dtypes(exclude=['object']).columns)
```

## 5 Go through


Let's replace NaNs with something first.

```
In [65]: train.replace('NaN', -999, inplace=True)
```

Let's calculate how many times one feature is greater than the other and create cross tabel out of it.

```
In [21]: # select first 42 numeric features
        feats = num_cols[:42]

        # build 'mean(featl > feat2)' plot
        gt_matrix(feats,16)
```



Indeed, we see interesting patterns here. There are blocks of features where one is strictly greater than the other. So we can hypothesize, that each column corresponds to cumulative counts, e.g. feature number one is counts in first month, second -- total count number in first two month and so on. So we immediately understand what features we should generate to make tree-based models more efficient: the differences between consecutive values.



## 5.1 VAR\_0002, VAR\_0003

```
In [59]: hist_it(train['VAR_0002'])
plt.ylim((0,0.05))
plt.xlim((-10,1010))

hist_it(train['VAR_0003'])
```

```
plt.ylim((0,0.03))
plt.xlim((-10,1010))
```

Out[59]: (-10, 1010)



```
In [39]: train['VAR_0002'].value_counts()
```

Out[39]:

12	5264
24	4763
36	3499
60	2899
6	2657
13	2478
72	2243
48	2222
3	2171
4	1917
2	1835
84	1801
120	1786
1	1724
7	1671
26	1637
5	1624

14	1572
18	1555
8	1513
999	1510
25	1504
96	1445
30	1438
9	1306
144	1283
15	1221
27	1186
38	1146
37	1078

...

877	1
785	1
750	1
653	1
784	1
764	1
751	1
797	1
926	1
691	1
808	1
774	1
902	1
755	1
656	1
814	1
813	1
685	1
739	1
935	1
906	1
807	1
550	1
933	1
804	1
675	1
674	1
745	1
778	1
851	1

Name: VAR\_0002, Length: 820, dtype: int64

In [38]: train['VAR\_0003'].value\_counts()

Out[38]: 0        17436



24	3469
12	3271
60	3054
36	2498
72	2081
48	2048
6	1993
1	1797
3	1679
84	1553
2	1459
999	1428
4	1419
120	1411
7	1356
13	1297
18	1296
96	1253
14	1228
8	1216
5	1189
9	1182
30	1100
25	1100
144	1090
15	1047
61	1008
26	929
42	921
...	
560	1
552	1
550	1
804	1
543	1
668	1
794	1
537	1
531	1
664	1
632	1
709	1
597	1
965	1
852	1
648	1
596	1
466	1

```

592      1
521      1
533      1
636      1
975      1
973      1
587      1
523      1
584      1
759      1
583      1
570      1
Name: VAR_0003, Length: 588, dtype: int64

```

We see there is something special about 12, 24 and so on, so we can create another feature  $x \bmod 12$ .


## 5.2 VAR\_0004

```

In [49]: train['VAR_0004_mod50'] = train['VAR_0004'] % 50
         hist_it(train['VAR_0004_mod50'])
         plt.ylim((0,0.6))

```

Out [49]: (0, 0.6)



## 6 Categorical features

Let's take a look at categorical features we have.

```

In [57]: train.loc[:,cat_cols].head().T

```

```

Out [57]:
          0          1  \
VAR_0001    H          H
VAR_0005    C          B
VAR_0008  False      False
VAR_0009  False      False

```

VAR_0010	False	False
VAR_0011	False	False
VAR_0012	False	False
VAR_0043	False	False
VAR_0044	[]	[]
VAR_0073	NaT	2012-09-04 00:00:00
VAR_0075	2011-11-08 00:00:00	2011-11-10 00:00:00
VAR_0156	NaT	NaT
VAR_0157	NaT	NaT
VAR_0158	NaT	NaT
VAR_0159	NaT	NaT
VAR_0166	NaT	NaT
VAR_0167	NaT	NaT
VAR_0168	NaT	NaT
VAR_0169	NaT	NaT
VAR_0176	NaT	NaT
VAR_0177	NaT	NaT
VAR_0178	NaT	NaT
VAR_0179	NaT	NaT
VAR_0196	False	False
VAR_0200	FT LAUDERDALE	SANTEE
VAR_0202	BatchInquiry	BatchInquiry
VAR_0204	2014-01-29 21:16:00	2014-02-01 00:11:00
VAR_0214	NaN	NaN
VAR_0216	DS	DS
VAR_0217	2011-11-08 02:00:00	2012-10-02 02:00:00
VAR_0222	C6	C6
VAR_0226	False	False
VAR_0229	False	False
VAR_0230	False	False
VAR_0232	True	False
VAR_0236	True	True
VAR_0237	FL	CA
VAR_0239	False	False
VAR_0274	FL	MI
VAR_0283	S	S
VAR_0305	S	S
VAR_0325	-1	H
VAR_0342	CF	EC
VAR_0352	0	0
VAR_0353	U	R
VAR_0354	0	R
VAR_0404	CHIEF EXECUTIVE OFFICER	-1
VAR_0466	-1	I
VAR_0467	-1	Discharged
VAR_0493	COMMUNITY ASSOCIATION MANAGER	-1
VAR_1934	IAPS	IAPS

	2	3	4
VAR_0001	H	H	R
VAR_0005	C	C	N
VAR_0008	False	False	False
VAR_0009	False	False	False
VAR_0010	False	False	False
VAR_0011	False	False	False
VAR_0012	False	False	False
VAR_0043	False	False	False
VAR_0044	[]	[]	[]
VAR_0073	NaT	NaT	NaT
VAR_0075	2011-12-13 00:00:00	2010-09-23 00:00:00	2011-10-15 00:00:00
VAR_0156	NaT	NaT	NaT
VAR_0157	NaT	NaT	NaT
VAR_0158	NaT	NaT	NaT
VAR_0159	NaT	NaT	NaT
VAR_0166	NaT	NaT	NaT
VAR_0167	NaT	NaT	NaT
VAR_0168	NaT	NaT	NaT
VAR_0169	NaT	NaT	NaT
VAR_0176	NaT	NaT	NaT
VAR_0177	NaT	NaT	NaT
VAR_0178	NaT	NaT	NaT
VAR_0179	NaT	NaT	NaT
VAR_0196	False	False	False
VAR_0200	REEDSVILLE	LIBERTY	FRANKFORT
VAR_0202	BatchInquiry	BatchInquiry	BatchInquiry
VAR_0204	2014-01-30 15:11:00	2014-02-01 00:07:00	2014-01-29 19:31:00
VAR_0214	NaN	NaN	NaN
VAR_0216	DS	DS	DS
VAR_0217	2011-12-13 02:00:00	2012-11-01 02:00:00	2011-10-15 02:00:00
VAR_0222	C6	C6	C6
VAR_0226	False	False	False
VAR_0229	False	False	False
VAR_0230	False	False	False
VAR_0232	True	False	True
VAR_0236	True	True	True
VAR_0237	WV	TX	IL
VAR_0239	False	False	False
VAR_0274	WV	TX	IL
VAR_0283	S	S	S
VAR_0305	P	P	P
VAR_0325	R	H	S
VAR_0342	UU	-1	-1
VAR_0352	R	R	R
VAR_0353	R	R	U
VAR_0354	-1	-1	0
VAR_0404	-1	-1	-1

VAR_0466	-1	-1	-1
VAR_0467	-1	-1	-1
VAR_0493	-1	-1	-1
VAR_1934	IAPS	RCC	BRANCH

VAR\_0200, VAR\_0237, VAR\_0274 look like some geographical data thus one could generate geography related features, we will talk later in the course.

There are some features, that are hard to identify, but look, there a date columns VAR\_0073 -- VAR\_0179, VAR\_0204, VAR\_0217. It is useful to plot one date against another to find relationships.


```
In [74]: date_cols = [u'VAR_0073', 'VAR_0075',
                      u'VAR_0156', u'VAR_0157', u'VAR_0158', 'VAR_0159',
                      u'VAR_0166', u'VAR_0167', u'VAR_0168', u'VAR_0169',
                      u'VAR_0176', u'VAR_0177', u'VAR_0178', u'VAR_0179',
                      u'VAR_0204',
                      u'VAR_0217']

for c in date_cols:
    train[c] = pd.to_datetime(train[c], format = '%d%b%y:%H:%M:%S')
    test[c] = pd.to_datetime(test[c], format = '%d%b%y:%H:%M:%S')

In [75]: c1 = 'VAR_0217'
        c2 = 'VAR_0073'

# mask = (~test[c1].isnull()) & (~test[c2].isnull())
# sc2(test.ix[mask, c1].values, test.ix[mask, c2].values, alpha=0.7, c = 'black')

mask = (~train[c1].isnull()) & (~train[c2].isnull())
sc2(train.loc[mask, c1].values, train.loc[mask, c2].values, c=train.loc[mask, 'target'].valu
```



We see that one date is strictly greater than the other, so the difference between them can be a good feature. Also look at horizontal line there -- it also looks like NaN, so I would rather create a new binary feature which will serve as an indicator that our time feature is NaN.